

Query Optimization Example

- Sailors (sid, sname, rating, age)
- Boats(bid, bname, color)
- Reserves(sid, bid, day, rname)
- Query:

```
SELECT S.sid, S.sname, S.age
FROM   Sailors S, Boats B, Reserves R
WHERE  B.bid = R.rid AND B.bid = R.bid AND
       B.color = "Red" AND S.age < 30;
```
- Reserves has 1000 pages, 10 tuples/page
- Sailors has 500 pages, 50 tuples/page
- Boats has 160 pages, 10 tuples/page
- Data is evenly distributed (assumption)

CPSC 421, 2009

1

Steps

- Query optimization steps:
 - Translate SQL Query to Relational Algebra Query
 - Create a query tree
 - Create left-deep alternative trees
 - Create query plans for our trees
 - Estimate costs of plans
 - Pick best one

CPSC 421, 2009

2

Step 1

- Translate SQL Query to Relational Algebra Query


$$\pi_{\text{sid, sname, age}}(\sigma_{\text{age} < 30}(\sigma_{\text{color} = \text{"Red"}}(\sigma_{\text{bid} = \text{bid}}(\mathbf{B} \times \sigma_{\text{sid} = \text{sid}}(\mathbf{S} \times \mathbf{R}))))))$$

CPSC 421, 2009

3

Step 2

- Create a query tree



Note our query only involves joins
(as opposed to plain old cross-
products) ...

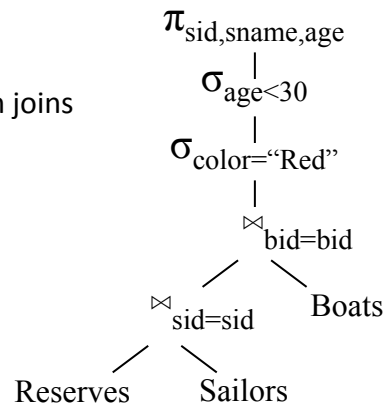
Lets draw the trees with joins

CPSC 421, 2009

4

Step 3

- Create left-deep alternative trees
 - We replace cross products if they are really joins
 - How?
 - Using RA equivalences!
 - This is a left-deep plan with joins





CPSC 421, 2009

5

Step 3 – More Alternatives

- Would we expect these to have different costs over previous alternative?



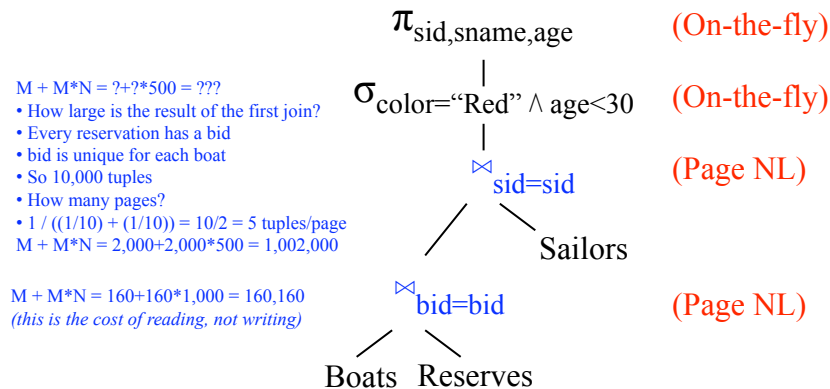
CPSC 421, 2009

6

Step 4 & 5 – Plans and Estimates

- Lets create some plans and estimates

Total cost: $160,160 + 1,002,000 = 1,162,160$ I/Os !!!



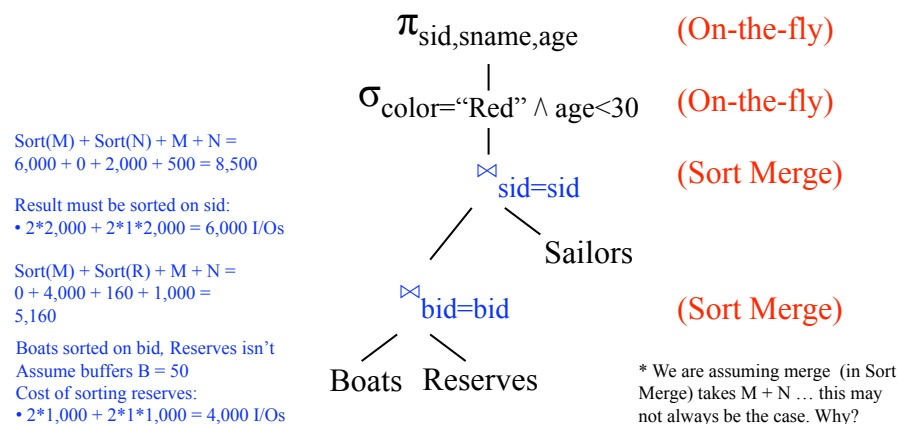
CPSC 421, 2009

7

Step 4 & 5 – Plans and Estimates

- Assume Boats has a clustered B+ Tree on bid (key)
- Assume S has a clustered B+ Tree on sid (key)

Total cost: $5,160 + 8,500 = 13,660$ I/Os !!!



CPSC 421, 2009

8

Note on Sort Merge

This is a much better plan than our previous page nested loops one

- But it is an overestimate for Sort-Merge join!
 - In Pass 0 we read all pages then write all pages ($2 \cdot M$)
 - We do not need to read all pages if they are pipelined to the next join ($1 \cdot M$)
 - In last merge pass, we don't have to write the last set of pages since we always pipeline them to the next operator
 - We also don't have to read in the file again:
 - $\text{Sort}(M) + \text{Sort}(N)$ if N needs to be sorted
 - $\text{Sort}(M) + N$ if N is sorted

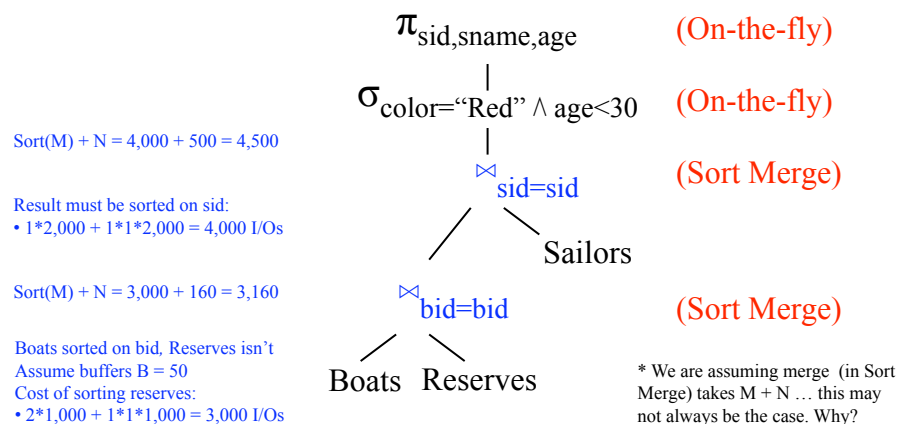
CPSC 421, 2009

9

Step 4 & 5 – Plans and Estimates

- Assume Boats has a clustered B+ Tree on bid (key)
- Assume S has a clustered B+ Tree on sid (key)

Total cost: $3,160 + 4,500 = 7,660$ I/Os !!!



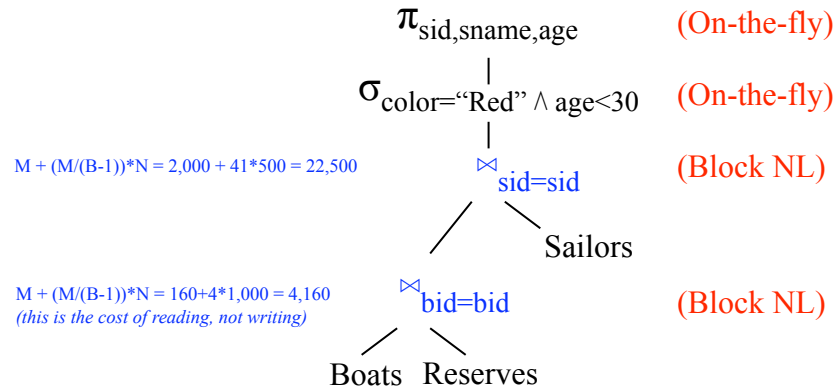
CPSC 421, 2009

10

Step 4 & 5 – Plans and Estimates

- Lets try block nested loop join (B=50)

Total cost: $4,160 + 22,500 = 26,660$ I/Os !!!



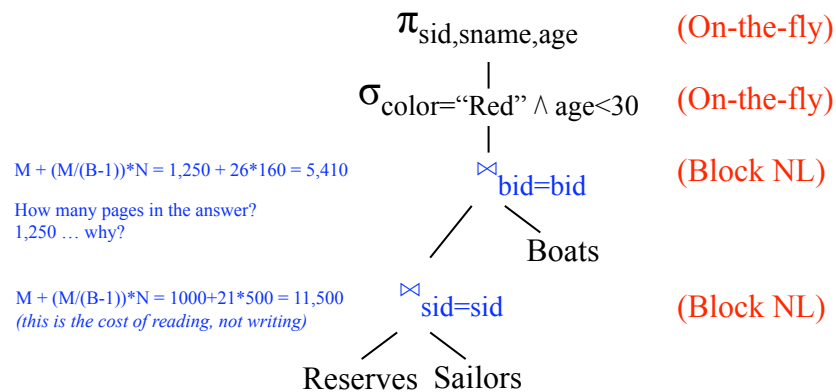
CPSC 421, 2009

11

Step 4 & 5 – Plans and Estimates

- Can we do better?

Total cost: $5,410 + 11,500 = 16,910$ I/Os !!!



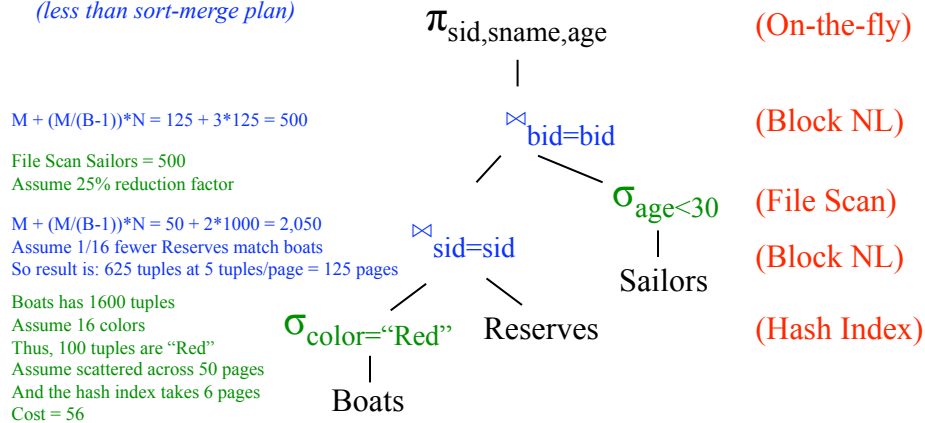
CPSC 421, 2009

12

Step 4 & 5 – Plans and Estimates

- Can we do even better?
 - Assume Boats has a hash index on color

Total cost: $56 + 2,050 + 500 + 500 = 3,106$ I/Os !!!
(less than sort-merge plan)



CPSC 421, 2009

13

Step 4 & 5 – Plans and Estimates

- Can we still do better?
 - Yes!
 - E.g., we could project on bid after first join ...
- Notice that by adding an index on color to Boats ...
 - We reduced our best query time in half!
 - This is one reason we talk about query optimization
 - It drives physical database design
 - Improving performance (by adding indexes, e.g.) should be driving on how query optimization works!

CPSC 421, 2009

14